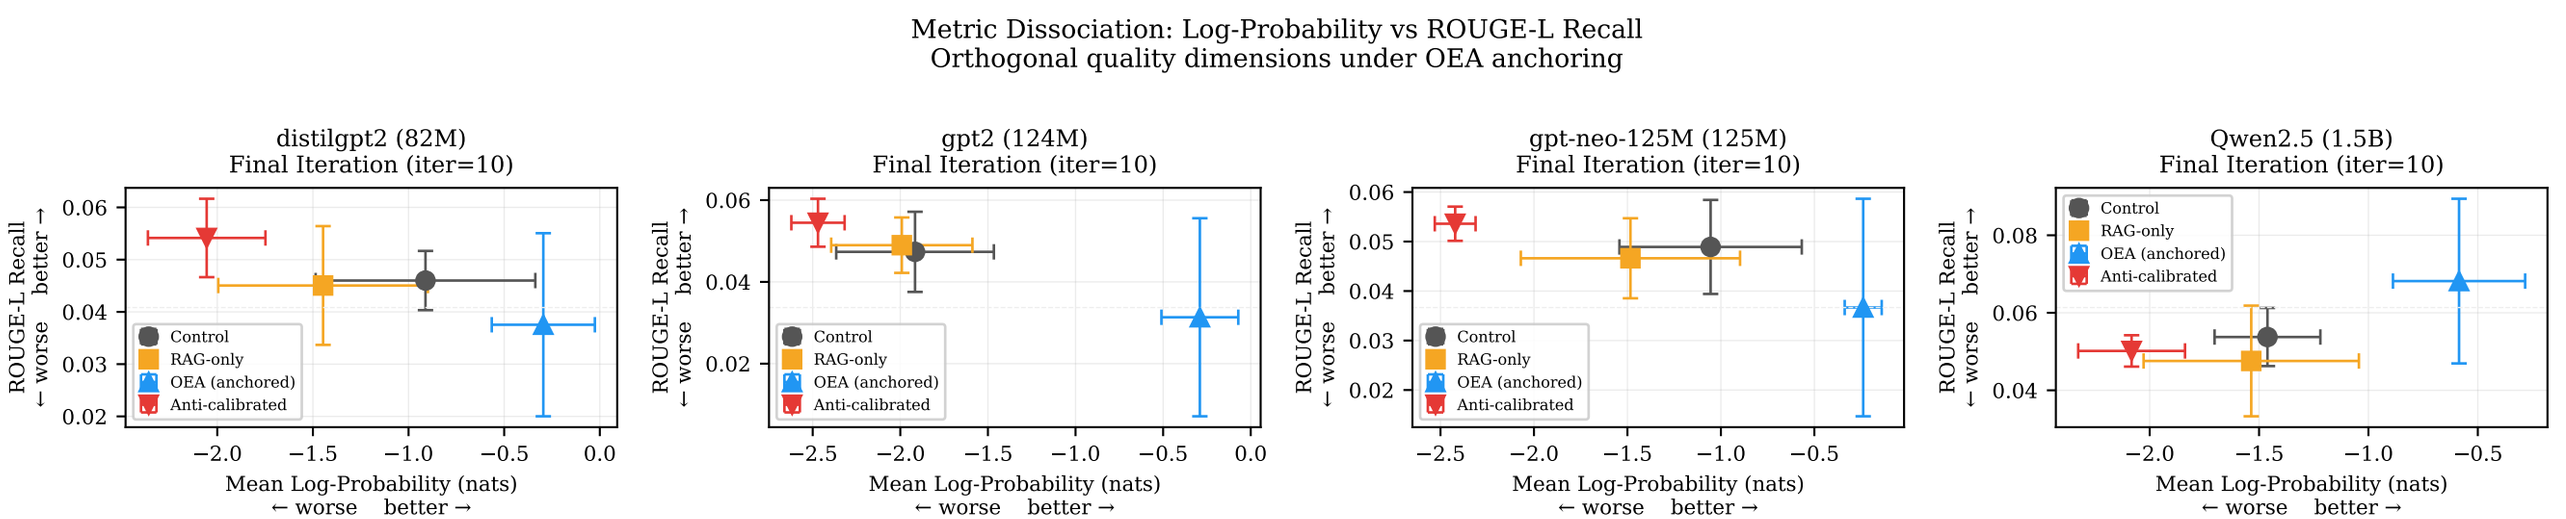
Metric Dissociation: Log-Probability vs ROUGE-L Recall
Orthogonal quality dimensions under OEA anchoring
distilgpt2 (82M)
Final Iteration (iter=10)
gpt2 (124M)
Final Iteration (iter=10)
gpt-neo-125M (125M)
Final Iteration (iter=10)
Qwen2.5 (1.5B)
Final Iteration (iter=10)
ROUGE-L Recall
← worse better →
Mean Log-Probability (nats)
← worse better →
Control
RAG-only
OEA (anchored)
Anti-calibrated