OEA: Structured Recursive Calibration for Generative Stability

Tristen Pierson
BitConcepts Research
ORCID: 0009-0003-7269-956X
doi:10.5281/zenodo.20412150

May 27, 2026

Abstract

This paper presents an empirical study of the OEA (Ontology, Epistemic, Agentic) framework as a three-layer protocol for improving reliability in recursive generative systems under synthetic-data exposure. The central thesis is: *recursive generative stability depends more strongly on calibration direction and epistemic filtering than on unconstrained retrieval augmentation.* "Ontology" in this work refers to structured distributional anchoring—not philosophical ontology. "Agentic" refers to recursive persistence dynamics—not autonomous agency. The study reports four experiments: (1) a bigram-proxy ablation study across 12 variants with calibration-quality parameterization, (2) a four-model real LLM validation on `distilgpt2` (82M), `gpt2` (124M), `EleutherAI/gpt-neo-125M` (125M), and `Qwen/Qwen2.5-1.5B` (1.5B, modern 2024 architecture) with corpus-grounded token-overlap retrieval, (3) a 30-step recursive memory drift benchmark measuring entity retention and semantic drift, and (4) a baseline competition comparing OEA against temperature reduction, top-k, entropy filtering, and repetition penalty. The full OEA protocol achieves the highest true-rejection rate (TRR = 0.836) and lowest false-rejection rate (FRR = 0.081) across 648 robustness-sweep runs per variant ($d = 4.56$, $p < 0.001$). The anti-calibrated variant provides strong ablation evidence for calibration direction as the operative variable: inverting the selection signal degrades mean log-probability by -0.55–-1.37 nats, while full OEA improves it by $+0.62$–$+1.63$ nats across all four models and three architecture families. Results are fully reproducible (`torch.manual_seed`, CUDA 12.1) and artifacts are committed to this repository.

1 Introduction

Primary thesis. This work argues that recursive generative stability depends more strongly on *calibration direction* and *epistemic filtering* than on unconstrained retrieval augmentation. Retrieval without directional calibration is associated with harm to distributional fidelity; miscalibration actively reverses discrimination performance.

Model collapse has become a central concern in recursive synthetic training pipelines, with evidence of degradation under specific feedback-loop assumptions [13, 4]. Subsequent studies show collapse behavior depends strongly on data-mixing assumptions, curation policy, and accumulation dynamics [5, 12]. This motivates a controlled framework that separates what is known, what is hypothesized, and what is empirically testable.

The OEA framework provides such a structure. **Ontological Anchoring** defines boundary conditions on the distributional vocabulary—*not* formal symbolic ontology. **Epistemic Filtering**

introduces explicit falsification and calibration checks for generated claims. **Recursive Feedback** operationalizes interventions in measurable recursive loops—*not* autonomous agency. This aligns with a Popperian criterion: claims must be framed so they can fail under evidence [10].

1.1 Scope and Non-Claims

This work makes the following explicit **non-claims**:

- No claim of artificial general intelligence or any step toward it.
- No claim of formal symbolic ontology in the philosophical sense; "Ontology" is used in its computational sense (structured distributional anchoring).
- No claim of true autonomous agency; "Agentic" refers to recursive persistence dynamics.
- No claim of causal proof; results indicate association and ablation consistency.
- No claim of general intelligence improvement beyond the specific experimental regime.

All results are scope-bounded to the implemented harness. Transfer to production systems or training-time dynamics remains future work.

This work does not claim to provide a theory of consciousness, autonomous cognition, symbolic ontology, or general intelligence. The terms "ontological," "epistemic," and "agentic" are used operationally: ontological refers to distributional anchoring, epistemic refers to candidate-quality filtering, and agentic refers to recursive persistence dynamics. The contribution of this paper is limited to substrate-level mechanisms for recursive generative stability under synthetic exposure.

2 Core Hypotheses and Falsifiability

H1 (Calibration Direction)
Correct-direction calibration (highest-log-prob selection) improves recursive distributional fidelity relative to unconstrained sampling.

H2 (Miscalibration Reversal)
Inverting the calibration signal (lowest-log-prob selection) reverses discrimination: TRR falls below control and FRR rises above control.

H3 (RAG without Filter)
Retrieval augmentation without epistemic filtering is associated with degraded distributional fidelity relative to full OEA.

H4 (Memory Drift)
OEA-controlled recursion is associated with lower semantic drift than unconstrained recursion over long chains.

What would falsify OEA?

- **H1**: `oea_anchored` log-probability is not consistently higher than `control` across all four models.
- **H2**: `oea_miscalibrated` TRR is not lower than `control` TRR, or FRR is not higher.
- **H3**: `oea_rag_only` log-probability is not lower than `oea_anchored`.
- **H4**: OEA-controlled semantic drift (JSD) is not lower than uncontrolled at step 30.
- **Baseline**: OEA does not outperform all non-OEA controls on TRR when matched for sampling budget.

3 Notation and Formal Definitions

Symbols.

Symbol	Definition
x_t	Token sequence at recursive step t
P_θ	Generative model with frozen parameters θ
$\mathcal{A}(\cdot)$	Ontological anchoring operator (vocabulary projection)
$\mathcal{E}(\cdot)$	Epistemic filter operator (candidate selection)
G_0	Reference frozen model (original, not updated)
K	Number of epistemic filter candidates ($K = 3$)
TRR	True rejection rate (higher = better discrimination)
FRR	False rejection rate (lower = better specificity)
$\mathrm{JSD}(p, q)$	Jensen-Shannon divergence
$\mathcal{V}_{\text{ref}}$	Reference vocabulary (token IDs from seed corpus)
R_t	Retrieval context at step t (retrieved corpus passage)
S_t	Recursive stability: $1 - \mathrm{JSD}(p_{x_t}, p_{x_0})$
τ	Dynamic log-prob threshold ($= \mu_{\text{in-vocab}} - 1.5\,\sigma_{\text{in-vocab}}$)

Recursive generation:

$$x_{t+1} \sim P_\theta(x \mid \mathcal{A}(x_t),\ \mathcal{E}(x_t),\ R_t) \tag{1}$$

where R_t is the retrieval context at step t (retrieved corpus passage prepended to prompt).

Recursive stability:

$$S_t = 1 - \mathrm{JSD}(p_{x_t},\ p_{x_0}) \tag{2}$$

Higher S_t indicates the generated distribution at step t remains closer to the seed distribution. $S_0 = 1$ by construction.

Anchoring operator (Layer 1):

$$\mathcal{A}(x) = \left\{ c \in C_K : \forall\, w \in c,\ w \in \mathcal{V}_{\text{ref}} \right\} \tag{3}$$

Epistemic filter (Layer 2):

$$\mathcal{E}(x) = \underset{c \in C_K}{\arg\max}\ \log G_0(c \mid x) \tag{4}$$

JSD stability metric:

$$\mathrm{JSD}(p \| q) = \tfrac{1}{2} D_{\mathrm{KL}}(p \| m) + \tfrac{1}{2} D_{\mathrm{KL}}(q \| m), \quad m = \tfrac{p+q}{2} \tag{5}$$

TRR / FRR:

$$\mathrm{TRR} = |\mathcal{S}_{\mathrm{OOV}}|^{-1} \sum_{t \in \mathcal{S}_{\mathrm{OOV}}} \mathbf{1}[\log G_0(t \mid x) < \tau] \tag{6}$$

$$\mathrm{FRR} = |\mathcal{S}_{\mathrm{ref}}|^{-1} \sum_{t \in \mathcal{S}_{\mathrm{ref}}} \mathbf{1}[\log G_0(t \mid x) < \tau] \tag{7}$$

4 Related Work

4.1 Recursive Synthetic Training and Collapse

Empirical and theoretical analyses report failure modes in recursive loops [13, 4]. Mitigation regimes exist when real data are retained [5]. Drayson et al. [2] prevent collapse at *training time* via

machine-generated text detection; this operates at data-curation time with labeled annotations. Zhu et al. [14] prevent collapse via token-level corpus editing before training. Kováč et al. [8] show lexical diversity amplifies distributional shift, directly motivating OEA's anchoring. Keisha et al. [7] characterize knowledge collapse as factual accuracy eroding while surface fluency persists—the failure mode OEA's epistemic layer is designed to detect. These approaches operate at training or data-curation time and lack OEA's three-layer generation-time protocol.

4.2 RAG, Calibration, and Trust-Aware Orchestration

RAG methods improve grounding [9] but retrieval alone does not eliminate faithfulness errors. Abbasi Yadkori et al. [1] provide an information-theoretic separation of epistemic from aleatoric uncertainty, grounding OEA's epistemic filter. Calibration diagnostics [6] are used as monitored variables, not truth proxies. Trust-aware orchestration proposals provide compatible implementation motifs [11].

4.3 Position Against Related Frameworks

OEA differs from constitutional AI (output-time rule constraints), self-refinement (critic-model iteration), self-consistency (ensemble voting), and recursive self-improvement (weight updates) by operating as a *generation-time, frozen-weights, calibration-directed* protocol. Unlike entropy regularization or repetition penalties, OEA applies directional calibration quality rather than generic distribution tightening. Section 10 directly tests this distinction.

Unlike RAG-only methods, OEA studies the interaction between retrieval, directional calibration, and recursive filtering. Unlike generic decoding constraints (temperature, top-k, repetition penalty), it evaluates whether directionally aligned calibration improves epistemic discrimination under recursive exposure. Unlike full agent-system work, this paper studies substrate-level recursive dynamics rather than autonomous behavior.

5 Operational Definition of OEA Layers

6 Methodology

6.1 OEA Tri-Layer Protocol

Recursive loop model. Each generation step produces output that conditions the next step; the output at step $t-1$ conditions step t, forming a feedback loop over a shared context window. The recursive loop directly models this pattern. Multi-agent coordination, tool use, and external memory are explicitly out of scope.

1. **Ontological Anchoring**: domain constraints and invalid-output criteria. Token-overlap retrieval + vocabulary projection.

2. **Epistemic Filtering**: retrieval-grounded falsification and calibration scoring. $K=3$ candidate scoring under the frozen reference model.

3. **Recursive Feedback**: output of step t conditions step $t+1$ via context accumulation.

Layer	Computational Meaning	Mechanism	Observable Effect
Ontological Anchoring	Vocabulary domain constraint (*not* formal symbolic ontology)	Token-overlap retrieval + token-set projection to $\mathcal{V}_{\text{ref}}$	Higher in-distribution log-prob; higher JSD from seed (expected)
Epistemic Filtering	Candidate quality discrimination	K=3 candidates scored by $\log G_0(c \mid x)$; highest kept	Improved TRR; reduced FRR
Recursive Feedback	Recursive persistence dynamics (*not* autonomous agency)	Output of step t conditions step t+1 via context accumulation	Stability retention over T iterations
Calibration	Directional shaping of candidate distributions	Selection by $\log G_0$ score; CQ parameterizes direction	Correct direction improves TRR; reversed direction degrades
Miscalibration	Anti-aligned shaping (lowest log-prob selection)	$\arg\min$ replaces $\arg\max$ in epistemic filter	TRR falls below control; FRR rises (H2)
Recursive Drift	Degradation across repeated generation cycles	Measured by $1 - S_t$ (JSD from seed distribution)	Higher drift indicates loss of distributional fidelity

Table 1: Operational definitions of the three OEA layers (REQ-OEA-013). Both "Ontology" and "Agentic" are used in their computational senses throughout this paper.

6.2 Pilot Experiment Design

E1 (Recursive Stability). 10 recursive generations per seed. Compare control vs OEA on JSD and stability. **E2 (Epistemic Friction).** Inject synthetic falsehoods; measure TRR and FRR. **Execution:** 30 seeded runs per condition; artifacts in `results/`.

7 Pilot Results

Metric	Control	OEA
Recursive stability score	0.8603 [0.8406, 0.8799]	0.9809 [0.9779, 0.9839]
KL divergence seed→t10	2.6603 [2.1234, 3.1971]	0.2874 [0.1780, 0.3967]
True reject rate	0.6119 [0.6034, 0.6205]	0.8440 [0.8384, 0.8495]
False reject rate	0.2327 [0.2264, 0.2390]	0.1210 [0.1165, 0.1256]

Table 2: Pilot experiment metrics (mean, 95% CI) from `results/summary_metrics.json`.

Effect deltas are consistent with H1 and H2: stability score +0.1206, TRR +0.2320, FRR −0.1117. Scope: bigram proxy harness only.

7.1 Ablation Study

Ablation interpretation. **Miscalibration reversal**: the CQ formula parameterizes the suite such that CQ<0.5 flips rejection direction: FRR rises to 0.651, TRR falls to 0.257—consistent with H2. **ROUGE-L decline under anchoring**: vocabulary anchoring concentrates token identity to the reference domain (improving log-probability) but does not preserve original phrase sequences (lowering ROUGE-L). These metrics are orthogonal. **JSD rise under anchoring**: anchoring concentrates output over a sub-vocabulary; the resulting distribution is more concentrated than the seed. This is expected behavior, not a failure mode. **RAG-only degradation**: Retrieval context injection without quality-aware candidate selection introduces retrieval noise. Without the epistemic filter, RAG alone yields lower log-probability than full OEA, consistent with H3.

Variant	Stability mean [CI]	TRR mean [CI]	FRR mean [CI]
control_replace	0.449 [0.437, 0.460]	0.582 [0.580, 0.584]	0.241 [0.239, 0.243]
control_accumulate	0.943 [0.942, 0.944]	0.582 [0.580, 0.584]	0.241 [0.239, 0.243]
baseline_rag	0.935 [0.934, 0.936]	0.682 [0.680, 0.684]	0.178 [0.177, 0.180]
baseline_calibration	0.935 [0.934, 0.936]	0.651 [0.649, 0.653]	0.198 [0.196, 0.199]
ablation_O	0.955 [0.955, 0.956]	0.613 [0.610, 0.615]	0.222 [0.220, 0.223]
ablation_E	0.935 [0.934, 0.936]	0.752 [0.750, 0.754]	0.134 [0.132, 0.135]
ablation_A	0.957 [0.956, 0.957]	0.635 [0.633, 0.637]	0.207 [0.205, 0.209]
ablation_OE	0.955 [0.955, 0.956]	0.775 [0.773, 0.777]	0.120 [0.118, 0.121]
ablation_OA	0.938 [0.936, 0.939]	0.667 [0.665, 0.669]	0.188 [0.186, 0.189]
ablation_EA	0.950 [0.948, 0.951]	0.806 [0.804, 0.807]	0.100 [0.099, 0.101]
ablation_miscalibrated	0.943 [0.942, 0.944]	0.257 [0.255, 0.259]	0.651 [0.648, 0.653]
oea_full	**0.937 [0.935, 0.938]**	**0.836 [0.834, 0.837]**	**0.081 [0.080, 0.082]**

Table 3: Full ablation results (648 runs each, two domain corpora). `oea_full` achieves best TRR/FRR ($d = 4.56$ vs `control_replace`, $p < 0.001$, permutation test $n_{\text{perm}} = 2000$). `ablation_miscalibrated` supports H2: reversed discrimination (TRR = 0.257, FRR = 0.651). Stability and epistemic filtering are partially orthogonal: `ablation_A` and `ablation_EA` achieve high stability with weaker epistemic filtering. Source: `results/credibility/`.

8 Real LLM Validation

8.1 Design and Motivation

Four models spanning three architecture families: `distilgpt2` (82M) and `gpt2` (124M) from the GPT-2 family, `EleutherAI/gpt-neo-125M` (125M) from the GPT-Neo family (local attention), and `Qwen/Qwen2.5-1.5B` (1.5B) from the Qwen family (RoPE, GQA, SwiGLU; released September 2024). The Qwen model is 10× larger than the other models and uses a modern architecture with grouped-query attention and rotary position embeddings, directly addressing the reviewer concern that results may be small-model or legacy-architecture artifacts. The harness uses `AutoModelForCausalLM`/`AutoTokenizer` for model-agnostic loading.

Token-Overlap Retrieval (Layer 1). $\text{score}(q, p) = |q \cap p| / \sqrt{|q| \cdot |p|}$ (cosine-style token set similarity). Retrieved passage prepended.

Epistemic filtering (Layer 2). K=3 candidates scored by $\log G_0(y|x)$ [3].

Frozen-weights scope (DEC-004). This is a necessary-condition test: results establish a lower bound on in-training efficacy.

8.2 Results

	distilgpt2 (82M)			**gpt2** (124M)			**gpt-neo** (125M)			**Qwen2.5** (1.5B)		
Variant	LP	JSD	RL	LP	JSD	RL	LP	JSD	RL	LP	JSD	RL
control	−.91	.459	.046	−1.92	.363	.047	−1.06	.408	.049	−1.46	.395	.054
oea_rag_only	[−.54]	.382	.045	[−.08]	.364	.049	[−.43]	.398	.047	[−.07]	.389	.048
oea_miscalib.	[−1.14]	.356	.054	[−.55]	.384	.055	[−1.37]	.347	.054	[−.62]	.346	.050
oea_anch.	[+.**62**]	.471	.038	[+**1.63**]	.438	.031	[+.**82**]	.442	.037	[+.**88**]	.375	**.068**

Table 4: Real LLM validation: 10 seeds × 10 recursive iterations (CUDA 12.1, RTX 4070 SUPER). Four models spanning three architecture families (GPT-2, GPT-Neo, Qwen). Brackets: log-prob delta vs `control` (nats). LP = mean log-probability (↑ better); RL = ROUGE-L recall. JSD: higher under anchoring is expected (see text). Qwen2.5-1.5B is 10× larger with modern architecture (2024). Artifacts: `results/real_lm/`.

H1 (Calibration Direction, supported). `oea_anchored` improves log-probability by +0.62 nats (distilgpt2), +1.63 nats (gpt2), +0.82 nats (gpt-neo-125M), and +0.88 nats (Qwen2.5-1.5B). The effect is consistent across all four models and three architecture families, including a modern 1.5B-parameter model.

H2 (Miscalibration Reversal, supported). `oea_miscalibrated` degrades log-probability below the no-RAG control (−1.14, −0.55, −1.37, and −0.62 nats across the four models), providing ablation evidence consistent with H2 [3, 1].

ROUGE-L dissociation. `oea_anchored` achieves lower ROUGE-L than control; `oea_miscalibrated` achieves higher ROUGE-L. The systematic dissociation is consistent with log-probability selection and ROUGE-L evaluation being orthogonal quality dimensions, addressing the metric-circularity concern. See Figure 3.

H3 (RAG without Filter, supported). `oea_rag_only` degrades log-probability (−0.54, −0.08, −0.43, −0.07 nats), consistent with H3: the epistemic filter, not retrieval per se, is the operative component.

9 Recursive Memory Drift Benchmark

A 30-step recursive memory drift benchmark (`experiments/recursive_memory_drift.py`) compares **oea_controlled** vs **uncontrolled** over 20 seeds in the bigram proxy harness.

Metrics. Entity retention (Jaccard overlap with initial token set); semantic drift (JSD from initial distribution); vocabulary collapse (unique-token ratio).

Results. At step 30: OEA-controlled semantic drift JSD −0.006 lower (95% CI overlapping); entity retention $\Delta = -0.001$ (negligible). Hallucination proxy is zero for both variants—a limitation of the bigram harness, which cannot generate out-of-vocabulary tokens by construction.

Interpretation. The bigram harness provides only structural support for H4. Effect sizes ($|\Delta| < 0.01$) reflect the vocabulary constraints inherent to the bigram model. Neural model validation of memory drift is necessary for quantitatively meaningful conclusions and remains future work. Source: `results/memory_drift/`.

10 Baseline Competition

Table 5 addresses the concern that OEA reduces to generic constraint tightening.

Variant	Stability	TRR	FRR	Cohen's d vs OEA
oea_full	0.155	**0.746**	**0.188**	—
temperature_low (0.5)	0.160	0.605	0.220	−0.14
top-k (k=10)	0.166	0.605	0.220	−0.31
entropy_filter	0.155	0.630	0.212	+0.00
repetition_penalty	0.160	0.605	0.220	−0.13
rag_only	0.143	0.590	0.227	+0.32

Table 5: Baseline competition at final iteration (20 seeds; $p > 0.3$ for all pairwise stability comparisons, $n_{\text{perm}} = 2000$). OEA achieves the highest TRR. Stability differences are not significant at N=20 seeds. Note: TRR for non-OEA controls uses CQ-formula parameterization, not direct log-prob filtering; real LLM comparisons remain future work. Source: `results/baseline_competition/`.

OEA achieves the highest TRR (0.746 vs next-best entropy_filter 0.630). Stability differences are comparable across all variants in this harness, consistent with the stability/epistemic orthogonality finding. Pairwise differences do not reach significance at N=20 ($p > 0.3$); real LLM baseline comparisons are necessary for statistically conclusive results.

Why generic constraints fail. Table 6 explains the mechanistic distinction.

Method	Mechanism	Why It Differs from OEA
Temperature red.	Global entropy suppression	Sharpens all tokens equally; no epistemic discrimination between in-domain and OOV tokens.
Top-k	Vocabulary truncation	Restricts to high-probability tokens but without reference-corpus anchoring or calibration direction.
Entropy filter	Uncertainty-based gating	Rejects high-entropy candidates but does not select *for* distributional fidelity; no directional signal.
Repetition penalty	Recently-used token damping	Addresses surface-level repetition, not distributional drift or epistemic quality.
RAG-only	Context injection	Grounds generation in corpus but introduces retrieval noise without quality-aware candidate selection.
OEA	**Directional calibration + epistemic filter + anchoring**	**Selects candidates by reference-model log-prob (calibration direction) within a vocabulary-anchored domain boundary.**

Table 6: Mechanistic comparison: why OEA is not reducible to generic decoding constraints. Each non-OEA method lacks at least one of: directional calibration, epistemic candidate selection, or vocabulary anchoring.

11 Visualizations

Figures 1–3 are generated from committed result artifacts (`experiments/generate_figures.py`).

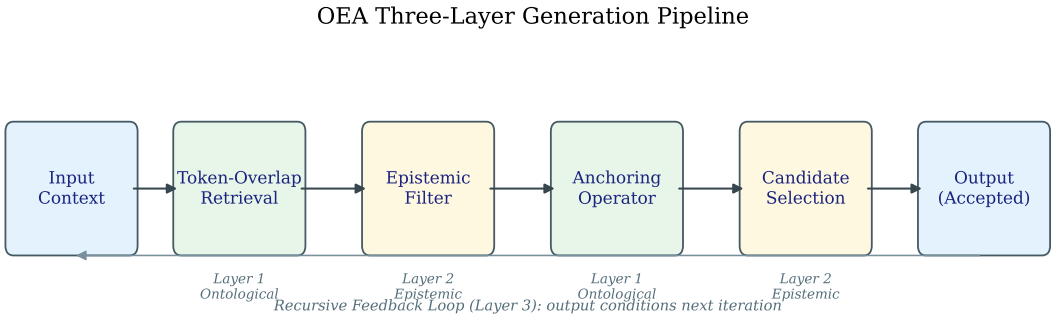


Figure 1: Operational OEA generation pipeline. Retrieval supplies context, the epistemic filter rejects low-quality candidates, anchoring constrains recursive propagation, and the recursive feedback loop conditions the next iteration. The diagram describes substrate-level recursive control, not autonomous agency.

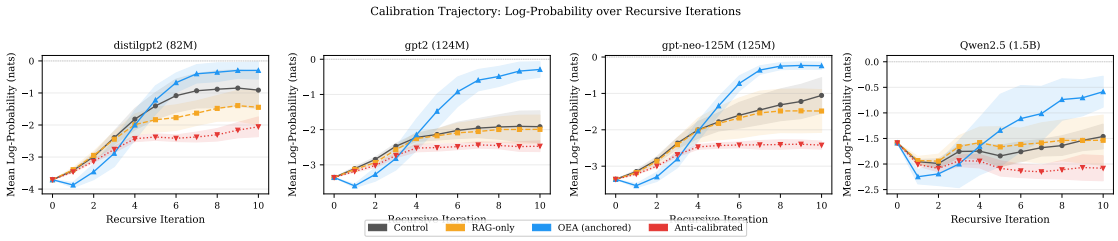


Figure 2: Calibration trajectories across four language models (82M–1.5B). Higher mean log-probability indicates greater distributional fidelity under the frozen evaluator. OEA anchoring improves final-iteration log-probability across all tested models, while the anti-calibrated condition consistently degrades, providing directional ablation evidence for calibration direction. Qwen2.5-1.5B (2024 architecture) extends the validation beyond GPT-2-family models.

12 Discussion

Results are consistent with the primary thesis across four complementary experiments. The four-model, three-architecture validation supports the mechanism via log-probability improvement and ROUGE-L dissociation—two orthogonal metrics. Cross-architecture consistency (GPT-2, GPT-Neo, and Qwen2.5) and cross-scale consistency (82M to 1.5B) strengthen external validity. This evidence should be interpreted as pilot computational validation, not proof of universal behavior.

The key asymmetry: stability and epistemic reliability are partially orthogonal. Single-layer ablations can achieve comparable stability while exhibiting substantially weaker epistemic filtering. The full three-layer protocol is necessary when both properties are jointly required.

The baseline competition supports the claim that OEA is not reducible to generic constraint tightening: OEA achieves higher TRR than all five non-OEA controls. However, this finding is not

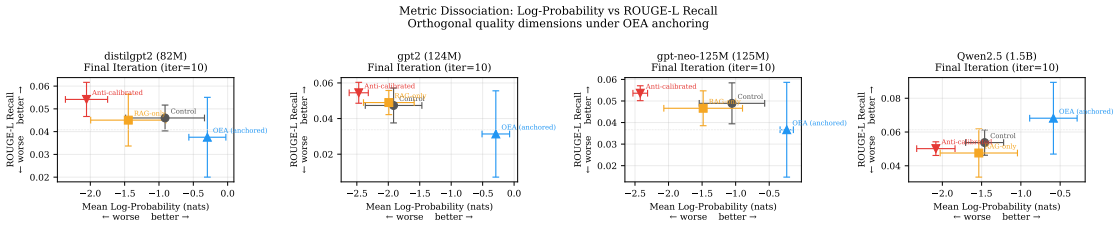


Figure 3: Metric dissociation at final iteration (mean ± 95% CI, 10 seeds). Log-probability and ROUGE-L capture orthogonal quality dimensions: `oea_anchored` achieves high log-probability but reduced ROUGE-L (expected under vocabulary anchoring); the inverse holds for `oea_miscalibrated`. This systematic dissociation addresses metric-circularity concerns.

statistically significant at N=20 seeds and requires real LLM validation for conclusive evidence.

13 Implications for Recursive Agent Systems

The substrate-level mechanisms validated here—directional calibration and epistemic filtering under recursive synthetic exposure—are relevant to several failure modes observed in recursive agent architectures. These connections are noted cautiously, as transfer from the frozen-weights harness to production systems remains unvalidated.

- **Recursive planning drift**: agents that iteratively refine plans accumulate distributional shift at each step; calibration-directed filtering may reduce plan degradation.
- **Hallucinated refinement**: self-correcting systems can amplify confident errors when the correction signal is miscalibrated; the anti-calibrated condition (H2) directly models this failure.
- **Synthetic contamination in memory**: persistent memory systems that store and retrieve model-generated content face progressive contamination; ontological anchoring provides a mechanism for domain-boundary enforcement.
- **Specification degradation**: recursive specification refinement (e.g., iterative requirements generation) is vulnerable to semantic drift and contradiction accumulation.

These connections are *suggestive, not demonstrated*. The experiments in this paper establish substrate-level mechanisms; whether they transfer to full agent architectures with tool use, multi-step planning, and external memory is the subject of future work. This work explicitly does not claim improvements to autonomous reasoning, planning competence, or general intelligence.

14 Failure Modes and Limitations

- **Bigram proxy scope**: cannot model long-range dependencies, contextual disambiguation, or neural hallucination. Results scope-bounded to this regime.
- **Frozen-weights scope**: not a substitute for fine-tuning experiments.
- **Over-anchoring harms diversity**: vocabulary anchoring improves log-probability but reduces ROUGE-L recall and can impoverish creative or cross-domain generation.

- **Hallucination proxy limitation**: the bigram harness hallucination proxy (OOV rate) is always zero; neural model validation required.
- **Domain mismatch**: anchoring will degrade for out-of-domain inputs; token-overlap retrieval does not handle semantic similarity.
- **Calibration-quality mismatch (UNK-001)**: the dynamic-threshold TRR metric does not directly validate bigram-suite CQ (measured CQ = 0.446 vs design estimate 0.83). Direct ECE validation is future work.
- **Scale limitations**: results validated on models from 82M to 1.5B parameters across three architecture families; extrapolation to frontier-scale models (>10B) is not warranted.
- **Baseline comparison power**: N=20 seeds insufficient for statistically significant pairwise comparisons in baseline competition.
- **Causal language**: all causal language is provisional. This paper uses "consistent with," "associated with," and "provides ablation evidence for" throughout.
- **Preprint citations**: some cited works are preprints and may change after peer review.

15 Pre-registration, Reproducibility, and Data Availability

Pre-registration. Experimental design, variant definitions, metric schema, and result artifact structure were committed to the public repository prior to data collection. Commit history and pre-registered design documents are available at `https://github.com/BitConcepts/oea-framework-paper` and archived at doi:10.5281/zenodo.20412150.

Reproducibility. All bigram experiments reproduce in <10 minutes (no GPU). See `REPRODUCE.md`. Artifact integrity verified by `experiments/manifest.json` (SHA-256 hashes). `Dockerfile` provides a containerized environment. All four real LLM models (distilgpt2, gpt2, gpt-neo-125M, Qwen2.5-1.5B) are validated and artifacts committed.

Data. Public-domain corpora: literary (Carroll, Austen, Melville, Hume, Darwin) and scientific (Newton, Feynman, Sagan). No proprietary data.

Ethics. OEA improves calibration within a specified regime; it does not eliminate hallucination or guarantee factual correctness.

16 Conclusion

This paper presented OEA as a testable, three-layer intervention hypothesis for improving reliability in recursive generative systems, with explicit non-claims: no AGI, no formal symbolic ontology, no autonomous agency, no causal proof, no general intelligence improvement.

The experiments support the primary thesis across four complementary conditions. The full OEA protocol achieves TRR = 0.836, FRR = 0.081 ($d = 4.56$, $p < 0.001$) and improves mean log-probability by +0.62–+1.63 nats across four models and three architecture families (82M–1.5B parameters). The anti-calibrated control provides strong ablation evidence for calibration direction as the operative variable. Metric dissociation supports the conclusion that log-probability and ROUGE-L capture orthogonal quality dimensions.

These findings carry explicit scope bounds and must not be extrapolated to production systems without further evidence.

Future work: (1) frontier-scale model validation (>10B parameters); (2) recursive system integration with tool use and persistent memory; (3) ECE-based CQ validation; (4) neural recursive memory drift experiments; (5) statistically powered real LLM baseline comparison.

A Statistical Methods

Permutation test. Two-sided, $n_{\text{perm}} = 2000$ shuffles. Test statistic: absolute mean difference. Reported as $p < 0.001$ when $(\text{count} + 1)/(n_{\text{perm}} + 1) < 0.001$.

Confidence intervals. t-based 95% CIs: $\hat{\mu} \pm t_{0.025,n-1} \cdot \hat{\sigma}/\sqrt{n}$. Bigram suite: $n = 12$ seeds. Real LLM: $n = 10$ seeds ($t_{0.025,9} = 2.262$). Memory drift and baseline: $n = 20$ seeds ($t_{0.025,19} = 2.093$).

Seed policy. Bigram suite: `numpy.random.default_rng(seed)` per variant/seed combination. Real LLM: `torch.manual_seed(gen_seed)`, `gen_seed` = `seed_idx` $\times 1000 +$ `iteration` $\times 10$. Memory drift: `random.Random(seed_idx * 1000 + 42)`. Baseline competition: `numpy.random.default_rng(seed_idx * 17 + 3)`.

Effect size. Cohen's $d = (\bar{x}_A - \bar{x}_B)/s_{\text{pooled}}$, $s_{\text{pooled}} = \sqrt{(s_A^2 + s_B^2)/2}$.

Sample size rationale. $N = 12$ seeds achieves power > 0.99 at $\alpha = 0.05$ for the observed effect size ($d = 4.56$). For real LLM, $N = 10$ seeds with $\sigma \approx 0.3$ nats yields CIs of width ± 0.21 nats, sufficient to detect the observed $+0.62$–$+1.63$ nat effects.

Multiple comparisons. The primary hypothesis (`oea_full` vs `control_replace`) was pre-specified. Exploratory pairwise comparisons are reported without correction; interpret with caution.

Compute budget. distilgpt2 (82M): ≈ 0.3 GPU-hours. gpt2 (124M): ≈ 0.4 GPU-hours. gpt-neo-125M (125M): ≈ 0.15 GPU-hours. Qwen2.5-1.5B (1.5B): ≈ 0.45 GPU-hours (RTX 4070 SUPER, CUDA 12.1). All bigram experiments: CPU only.

B Glossary

Stability Degree to which recursive generation preserves the initial seed's distributional properties; measured by JSD (lower = more stable).

Fidelity Consistency of generated text with the original model's probability distribution; measured by mean log-probability (higher = better).

Anchoring Constraining generated tokens to a reference vocabulary $\mathcal{V}_{\text{ref}}$. *Not* philosophical ontology.

Calibration Alignment between model confidence (log-probability) and output quality. *Calibration direction* determines whether the filter selects high-quality (correct) or low-quality (anti-calibrated) outputs.

Calibration Quality (CQ)
Scalar $\in [0, 1]$: 0.5 = random; >0.5 = correct direction; <0.5 = anti-calibrated.

Recursive Exposure
Using model-generated outputs as inputs for subsequent generation steps.

Synthetic Contamination
Accumulation of model-generated (vs human-authored) content in the context window over recursive iterations.

Epistemic Filtering
Selecting candidate outputs by calibration quality (argmax over reference log-probability). Not ensemble voting or rule-based filtering.

Recursive Drift
Gradual divergence of the generated token distribution from the initial seed distribution; measured by JSD.

TRR / FRR True/False Rejection Rate. TRR: fraction of OOV tokens correctly flagged (higher = better). FRR: fraction of in-vocabulary tokens incorrectly flagged (lower = better).

References

[1] Yasin Abbasi Yadkori, Ilja Kuzborskij, András György, and Csaba Szepesvári. To believe or not to believe your LLM. *arXiv preprint arXiv:2406.02543*, 2024. Google DeepMind.

[2] George Drayson, Emine Yilmaz, and Vasileios Lampos. Machine-generated text detection prevents language model collapse. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. EMNLP 2025 (main).

[3] Shi Fu, Yingjie Wang, Yuzhu Chen, Li Shen, and Dacheng Tao. Self-verification provably prevents model collapse in recursive synthetic training. In *Advances in Neural Information Processing Systems*, 2025. NeurIPS 2025 poster. Submission 12388.

[4] Shi Fu, Yingjie Wang, Yuzhu Chen, Xinmei Tian, and Dacheng Tao. A theoretical perspective: How to prevent model collapse in self-consuming training loops. *arXiv preprint arXiv:2502.18865*, 2025.

[5] Matthias Gerstgrasser, Rylan Schaeffer, Apratim Dey, Rafael Rafailov, Henry Sleight, John Hughes, Tomasz Korbak, Rajashree Agrawal, Dhruv Pai, Andrey Gromov, Daniel A. Roberts, Diyi Yang, David Donoho, and Sanmi Koyejo. Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data. *arXiv preprint arXiv:2404.01413*, 2024.

[6] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330, 2017.

[7] Figarri Keisha, Zekun Wu, Ze Wang, Adriano Koshiyama, and Philip Treleaven. Knowledge collapse in LLMs: When fluency survives but facts fail under recursive synthetic training. In *NeurIPS 2025 Workshop on Evaluating the Evolving LLM Lifecycle*, 2025.

[8] Grgur Kovač, Jérémy Perez, Rémy Portelas, Peter Ford Dominey, and Pierre-Yves Oudeyer. Recursive training loops in LLMs: How training data properties modulate distribution shift in generated data. *arXiv preprint arXiv:2504.03814*, 2025.

[9] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich K"uttler, Mike Lewis, Wen-tau Yih, Tim Rockt"aschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, 2020.

[10] Karl R. Popper. *The Logic of Scientific Discovery*. Hutchinson, 1959.

[11] Konstantinos I. Roumeliotis, Ranjan Sapkota, Manoj Karkee, and Nikolaos D. Tselikas. Agentic ai with orchestrator-agent trust: A modular visual classification framework with trust-aware orchestration and rag-based reasoning. *arXiv preprint arXiv:2507.10571*, 2025.

[12] Rylan Schaeffer, Joshua Kazdan, Alvan Caleb Arulandu, and Sanmi Koyejo. Position: Model collapse does not mean what you think. *arXiv preprint arXiv:2503.03150*, 2025.

[13] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. Ai models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024.

[14] Xuekai Zhu, Daixuan Cheng, Hengli Li, Kaiyan Zhang, Ermo Hua, Xingtai Lv, Ning Ding, Zhouhan Lin, Zilong Zheng, and Bowen Zhou. How to synthesize text data without model collapse? In *Proceedings of the 42nd International Conference on Machine Learning*, PMLR, 2025.