

OEA Framework

Full Experiment Insights Report

Structured Recursive Calibration

BitConcepts Research • 2026-05-13

1. Pilot Experiment — Recursive Stability & Epistemic Friction
 2. Credibility Suite — 12-Variant Ablation (7,776 runs)
3. Real LLM Validation — 3 Models, 2 Architecture Families
4. Agentic Memory Drift — 30-Step Recursive Benchmark
5. Baseline Competition — OEA vs 5 Non-OEA Controls
 6. Error Taxonomy & CQ Measurement

1. Pilot Experiment

Recursive stability & epistemic friction (30 seeds, bigram proxy)

Experiment Design

Seeds: 30 Iterations: 10 Token space: 12

Recursive Stability

Control KLD:	2.6603	[2.1234, 3.1971]
OEK KLD:	0.2874	[0.1780, 0.3967]
Control stab:	0.8603	[0.8406, 0.8799]
OEK stab:	0.9809	[0.9779, 0.9839]
Delta:	+0.1206	

Epistemic Friction

Control TRR:	0.6119	[0.6034, 0.6205]
OEK TRR:	0.8440	[0.8384, 0.8495]
Control FRR:	0.2327	[0.2264, 0.2390]
OEK FRR:	0.1210	[0.1165, 0.1256]
TRR delta:	+0.2320	
FRR delta:	-0.1117	

Interpretation

H1 supported: OEK stability +12.1% absolute over control.
H2 supported: TRR +23.2pp, FRR -11.2pp.
Scope: bigram proxy harness only.

2. Credibility Suite

12-variant ablation study (7,776 runs, $d=4.56$, $p<0.001$)

Configuration

Total runs: 7,776 (12 variants × 12 seeds × 3 depths × 3 ratios × 3 noise)
Corpora: public_domain_snippets, scientific_snippets

Headline: OEA Full vs Control (replace)

Stability delta: +0.4882 ($d = 4.56$, $p < 0.001$)
Perplexity delta: +231.9
TRR delta: +0.2541
FRR delta: -0.1600

Top 5 Variants by Stability

ablation_A	stab=0.9565	TRR=0.635	FRR=0.207
ablation_OE	stab=0.9554	TRR=0.775	FRR=0.120
ablation_0	stab=0.9550	TRR=0.613	FRR=0.222
ablation_EA	stab=0.9496	TRR=0.806	FRR=0.100
ablation_miscalibrated	stab=0.9433	TRR=0.257	FRR=0.651

Key Finding

oea_full: TRR=0.836, FRR=0.081 (best epistemic discrimination).
ablation_miscalibrated: TRR=0.257, FRR=0.651 (reversed – H2 confirmed).
Stability & epistemic filtering are partially orthogonal.

Credibility Suite — Aggregate Metrics

648 runs per variant across 2 domain corpora

Variant	Stability	TRR	FRR	Perplexity	N
ablation_A	0.9565	0.635	0.207	725.0	648
ablation_E	0.9347	0.752	0.134	452.7	648
ablation_EA	0.9496	0.806	0.100	436.4	648
ablation_O	0.9550	0.613	0.222	986.1	648
ablation_OA	0.9378	0.667	0.188	654.5	648
ablation_OE	0.9554	0.775	0.120	531.7	648
ablation_miscalibrate	0.9433	0.257	0.651	891.7	648
baseline_calibration_o	0.9347	0.651	0.198	452.7	648
baseline_rag_only	0.9347	0.682	0.178	452.7	648
control_accumulate	0.9433	0.582	0.241	891.7	648
control_replace	0.4485	0.582	0.241	270.2	648
oea_full	0.9367	0.836	0.081	502.2	648

3. Real LLM Validation

Three models, two architecture families (GPT-2, GPT-Neo)

Design

3 models × 4 variants × 10 seeds × 10 iterations = 1,200 steps

Hardware: RTX 4070 SUPER, CUDA 12.1

distilgpt2 (82M)

Control (iter 10): log_prob=-0.911 JSD=0.459 ROUGE-L=0.046

oea_anchored log_prob=+0.6160 JSD=-0.0118 ROUGE-L=-0.0085

oea_miscalibrated log_prob=-1.1439 JSD=+0.1027 ROUGE-L=0.0081

oea_rag_only log_prob=-0.5353 JSD=+0.0769 ROUGE-L=-0.0009

CQ measurement:

oea_rag_only: TRR=0.5625 → CQ=0.501

oea_anchored: TRR=0.5787 → CQ=0.520

oea_miscalibrated: TRR=0.6238 → CQ=0.571

gpt2 (124M)

Control (iter 10): log_prob=-1.916 JSD=0.363 ROUGE-L=0.0474

oea_anchored log_prob=+1.6255 JSD=-0.0742 ROUGE-L=-0.016

oea_miscalibrated log_prob=-0.5539 JSD=+0.0207 ROUGE-L=0.0071

oea_rag_only log_prob=-0.0755 JSD=+0.0006 ROUGE-L=0.0016

CQ measurement:

oea_rag_only: TRR=0.6312 → CQ=0.614

oea_anchored: TRR=0.4700 → CQ=0.445

oea_miscalibrated: TRR=0.6250 → CQ=0.607

gpt-neo-125M (125M)

Control (iter 10): log_prob=-1.055 JSD=0.408 ROUGE-L=0.0489

oea_anchored log_prob=+0.8163 JSD=-0.0341 ROUGE-L=-0.0122

oea_miscalibrated log_prob=-1.3664 JSD=+0.0610 ROUGE-L=0.0047

oea_rag_only log_prob=-0.4288 JSD=+0.0098 ROUGE-L=-0.0023

CQ measurement:

oea_rag_only: TRR=0.5475 → CQ=0.493

oea_anchored: TRR=0.4987 → CQ=0.438

oea_miscalibrated: TRR=0.5725 → CQ=0.521

Hypothesis Summary (all 3 models)

H1 (Calibration Direction): SUPPORTED (+0.62 to +1.63 nats)

H2 (Miscalibration Reversal): SUPPORTED (-0.55 to -1.37 nats)

H3 (RAG without Filter): SUPPORTED (-0.08 to -0.54 nats)

ROUGE-L dissociation: CONFIRMED across all 3 models

Real LLM — Final Iteration Results (iter 10)

10 seeds × 10 iterations per model/variant

Model	Variant	Log-prob	Δ Log-prob	JSD	ROUGE-L
distilgpt2	control	-0.911	—	0.459	0.0460
	oea_anchored	-0.295	+0.6160	0.471	0.0375
	oea_miscalibrated	-2.055	-1.1439	0.356	0.0542
	oea_rag_only	-1.447	-0.5353	0.382	0.0450
gpt2	control	-1.916	—	0.363	0.0474
	oea_anchored	-0.290	+1.6255	0.438	0.0313
	oea_miscalibrated	-2.470	-0.5539	0.343	0.0545
	oea_rag_only	-1.991	-0.0755	0.363	0.0490
gpt-neo-125M	control	-1.055	—	0.408	0.0489
	oea_anchored	-0.239	+0.8163	0.442	0.0367
	oea_miscalibrated	-2.421	-1.3664	0.347	0.0536
	oea_rag_only	-1.484	-0.4288	0.398	0.0466

4. Agentic Memory Drift Benchmark

30-step recursive summarization (20 seeds, bigram proxy)

Design

20 seeds × 2 variants × 30 steps
Bigram proxy recursive summarization

Final Step (step 30) Comparison

uncontrolled:		
entity_retention	0.1571	[0.1550, 0.1592]
semantic_drift	0.4937	[0.4823, 0.5051]
hallucination_proxy	0.0000	[0.0000, 0.0000]
vocab_collapse	0.9213	[0.9089, 0.9336]
oea_controlled:		
entity_retention	0.1561	[0.1535, 0.1587]
semantic_drift	0.4873	[0.4787, 0.4958]
hallucination_proxy	0.0000	[0.0000, 0.0000]
vocab_collapse	0.9150	[0.8999, 0.9301]

OEA vs Uncontrolled Deltas

entity_retention	-0.0011 ↓
semantic_drift	-0.0064 ↓
hallucination_proxy	+0.0000 ↓
vocab_collapse	-0.0063 ↓

Interpretation

Bigram harness provides structural support for H4 only.
Effect sizes $|\Delta| < 0.01$ reflect vocabulary constraints.
Hallucination proxy = 0.0 for both (bigram limitation).
Neural validation required for quantitative conclusions.

5. Baseline Competition

OEA vs 5 non-OEA controls (20 seeds, bigram proxy)

Design

20 seeds × 10 iterations
OEA vs 5 non-OEA controls: temperature_low, top-k, entropy_filter, repetition_penalty, rag_only.

Final Iteration Results

Variant	Stability	TRR	FRR	d(stab)	p(stab)
oea_full	0.1550	0.746	0.145	+0.000	1.000
temperature_low	0.1601	0.605	0.226	-0.135	0.668
top_k	0.1656	0.605	0.226	-0.312	0.307
entropy_filter	0.1550	0.630	0.211	+0.000	1.000
repetition_penalty	0.1597	0.605	0.226	-0.127	0.699
rag_only	0.1425	0.590	0.234	+0.317	0.328

Key Finding

OEA achieves highest TRR (0.746 vs next-best entropy_filter 0.630).
Stability differences not significant at N=20 (p > 0.3).
OEA is not reducible to generic constraint tightening.

6. Error Taxonomy

Mean error counts per variant (648 runs each)

Variant	Tail Loss	Novel Noise	False Accept	False Reject
ablation_A	0.0	337.6	109.5	62.1
ablation_E	0.0	0.0	74.3	40.1
ablation_EA	0.0	0.0	58.3	30.0
ablation_O	0.0	522.0	116.2	66.5
ablation_OA	0.0	164.2	100.0	56.3
ablation_OE	0.0	0.0	67.5	35.9
ablation_miscalibrated	0.0	522.0	223.0	195.2
baseline_calibration_only	0.0	0.0	104.7	59.3
baseline_rag_only	0.0	0.0	95.4	53.5
control_accumulate	0.0	522.0	125.5	72.3
control_replace	389.5	239.4	125.5	72.3
oea_full	0.0	0.0	49.2	24.3

CQ Measurement & Evidence Chain

Calibration quality measured from real LLM true-reject rates

CQ Measurement Summary (REQ-OEA-012)

Model	Variant	TRR	CQ
distilgpt2	oea_rag_only	0.5625	0.501
distilgpt2	oea_anchored	0.5787	0.520
distilgpt2	oea_miscalibrated	0.6238	0.571
gpt2	oea_rag_only	0.6312	0.614
gpt2	oea_anchored	0.4700	0.445
gpt2	oea_miscalibrated	0.6250	0.607
gpt-neo-125M	oea_rag_only	0.5475	0.493
gpt-neo-125M	oea_anchored	0.4987	0.438
gpt-neo-125M	oea_miscalibrated	0.5725	0.521

Evidence Chain Status

Design estimate CQ = 0.83 (credibility_suite.py).
Measured CQ for oea_anchored: 0.438–0.446 across models.
Mismatch documented as UNK-001. Direct ECE validation is future work.

Interpretation

Dynamic threshold adapts proportionally to anchoring,
so TRR does not improve despite better log-probability.
This is a known limitation, not a falsification of OEA.